

BIOC461

Research Design and Analysis in Biochemistry

Lecture 1

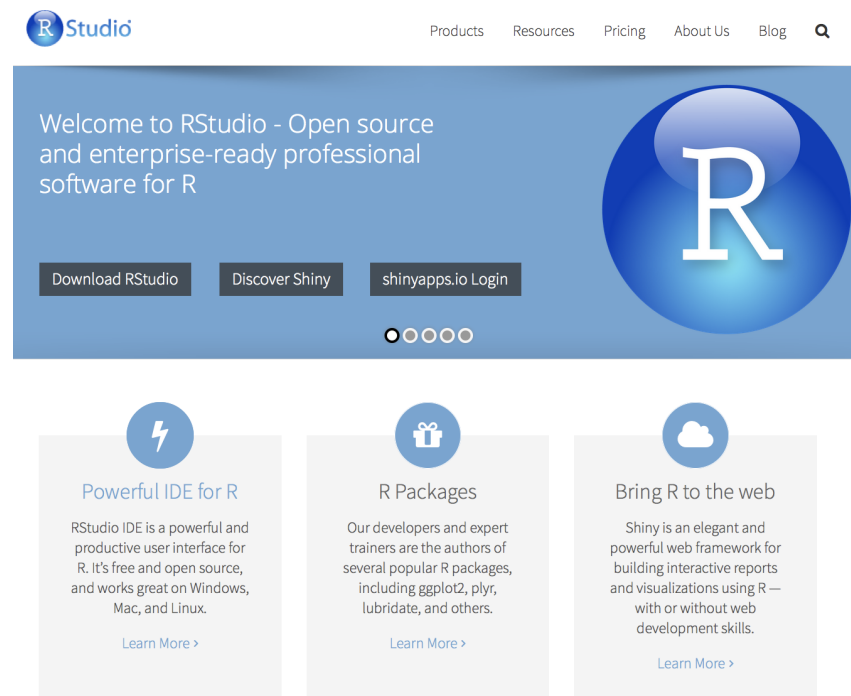
A/Prof Mik Black
Department of Biochemistry
19 March 2020

What's so great about statistics?

- In science we spend much of our time generating experimental data.
- We use this data both to answer our scientific questions, and to formulate new hypotheses.
- Rigorous analysis and interpretation of data is a vital component of science.
- Statistics provides a collection of tools to help us answer questions

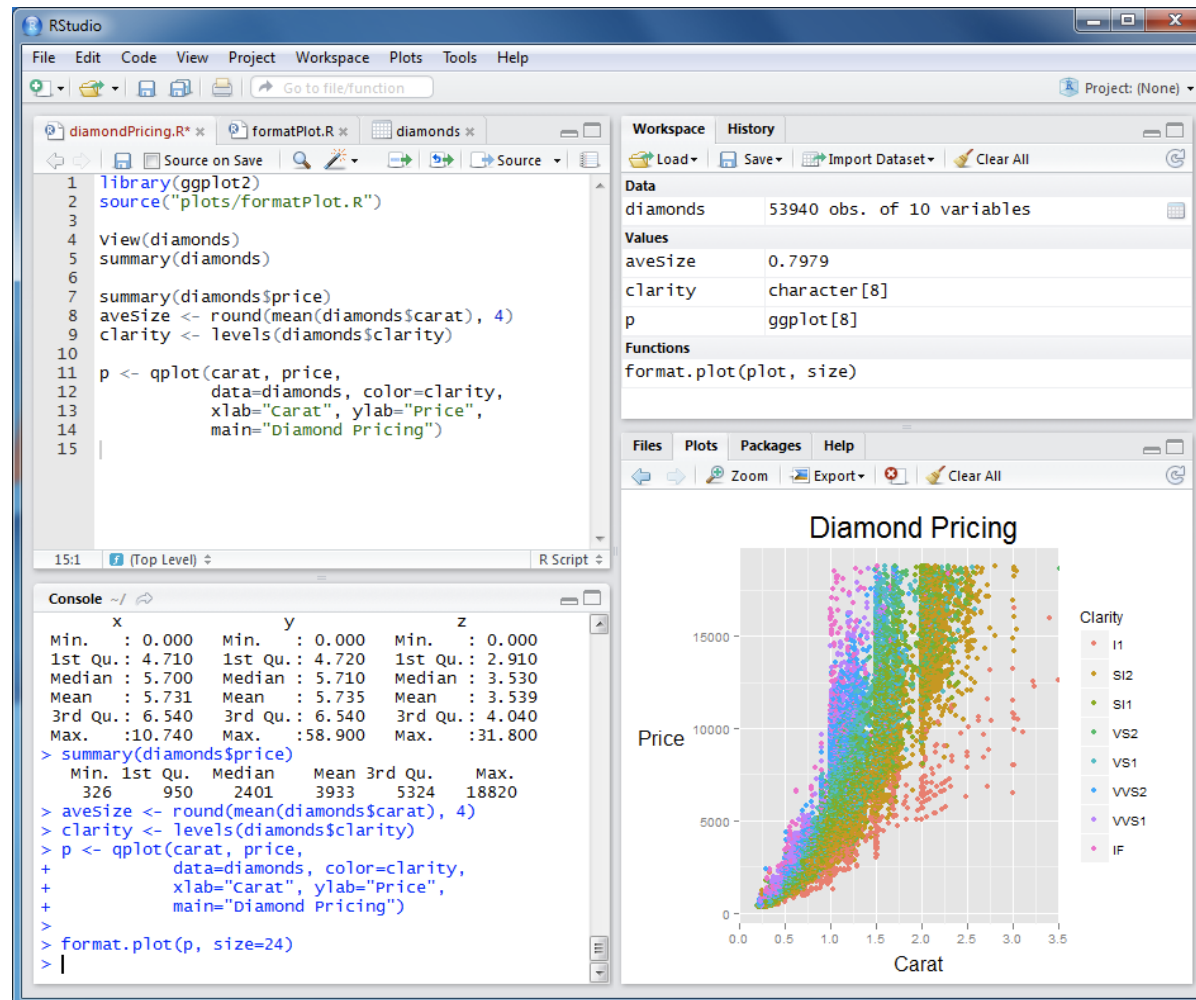
Introducing RStudio

The free RStudio application provides an “integrated development environment” (IDE) for working in R.



<https://www.rstudio.com/>

What does it look like?



Getting started with R

- Open RStudio and spend some time getting familiar with the layout:
 - Console: type commands here
 - Source: save your code here
 - Environment/History: session info
 - Files/Plots/Packages/Help/Viewer: everything else :)
- You can write commands/code in the source pane, and then “submit” it (i.e., run it) in the console by hitting `ctrl+Enter` (Windows) or `Cmd+Enter` (Mac)
- Comments start with `#` and are ignored by R

Getting started with R

```
# This is a comment
```

```
# This adds 3 + 5
```

```
3 + 5
```

```
## [1] 8
```

```
# This adds 3 + 2 and assigns it to an object called "a"
```

```
a = 3 + 2
```

```
# This shows the value of "a"
```

```
a
```

```
## [1] 5
```

Getting started with R

```
# This also adds 3 + 2 and assigns it to an object called "a"  
a <- 3 + 5
```

```
# This takes the object "a", and adds 10 to it  
a + 10
```

```
## [1] 18
```

```
# This calculates the square root of "a"  
# sqrt is called a "function" - we'll use functions a lot  
sqrt(a)
```

```
## [1] 2.828427
```

Lets do something useful

- While it is fun to use R as a giant calculator, that's not going to help (much) with your research.
- We want to do some data analysis in R
 - First, we need data
 - Good, clean, well-organised data

Data organization

- A spreadsheet application (e.g., Microsoft Excel) is often used to “organize” data.
 - Excel has the ability to make raw data look very pretty.
 - Statistics applications (e.g., SPSS, R, SAS, Prism) often have trouble reading in the pretty data.
- Try to keep it simple:
 - Samples in rows, variables in columns
 - No blank cells or rows (unless data are missing)

Example data:

Animal	Gender	Diet	<u>WeightGain_grams</u>
Mouse1	Male	HighFat	100
Mouse2	Female	HighFat	145
Mouse3	Female	HighFat	135
Mouse4	Male	LowFat	104
Mouse5	Female	LowFat	120
Mouse6	Male	LowFat	116
Mouse7	Male	LowFat	109

Note the lack of units in the Weight Gain column – stats programs expect numeric values for continuous data.

Loading the data

```
## Load the data into R from a CSV file  
mouseData <- read.csv('DataSets/mouseExample.csv')
```

```
## what sort of data object is it?  
class(mouseData)
```

```
## [1] "data.frame"
```

```
## What are its dimensions?  
dim(mouseData)
```

```
## [1] 7 4
```

Variable names

```
## List the variables names in the dataframe  
names(mouseData)
```

```
## [1] "Animal"          "Gender"          "Diet"  
## [4] "WeightGain_grams"
```

Inspecting the data

```
## Typing the name of the object displays its contents  
mouseData
```

```
##   Animal Gender    Diet WeightGain_grams  
## 1 Mouse1    Male HighFat             100  
## 2 Mouse2 Female HighFat             145  
## 3 Mouse3 Female HighFat             135  
## 4 Mouse4    Male  LowFat             104  
## 5 Mouse5 Female  LowFat             120  
## 6 Mouse6    Male  LowFat             116  
## 7 Mouse7    Male  LowFat             109
```

```
## Open the built-in RStudio viewer:  
View(mouseData)
```

Accessing individual variables

```
## The $ symbol allows you to reference the columns of a dataframe.  
## The following produces a summary of the data in the  
## "WeightGain_grams" column, and then calculates the standard  
## deviation.  
summary(mouseData$WeightGain_grams)
```

```
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.  
## 100.0   106.5   116.0   118.4   127.5   145.0
```

```
sd(mouseData$WeightGain_grams)
```

```
## [1] 16.46063
```

Accessing individual variables

You can also just use the column number
`summary(mouseData[,4])`

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    100.0   106.5   116.0   118.4   127.5   145.0
```

Or the name of the column
`summary(mouseData[, "WeightGain_grams"])`

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    100.0   106.5   116.0   118.4   127.5   145.0
```

Summarizing data

- Frequency Distribution: display the number of times each value occurs in the data set
- Example (for discrete quantitative data)
 - Size of litters for n=36 sows
 - Litter size (number of piglets) is an integer (discrete)

```
pigs<-read.csv( 'DataSets/pigs.csv' )  
head(pigs)
```

```
##  Weight  
## 1      10  
## 2      12  
## 3      10  
## 4       7  
## 5      14  
## 6      11
```


Frequency tables and histograms

- Summarize with a frequency distribution (as a table):

```
## Since there is only one variable, can just specify the  
## "pigs" data.frame  
table(pigs)
```

```
## pigs  
## 5 7 8 9 10 11 12 13 14  
## 1 2 3 3 9 8 5 3 2
```


```
## Better practice to specify the variable you are interested in  
table(pigs$Weight)
```

```
##  
## 5 7 8 9 10 11 12 13 14  
## 1 2 3 3 9 8 5 3 2
```

Frequency tables and histograms

- Generate a frequency histogram:

```
hist(pigs$Weight)
```



Example: heroin data

- A pathologist measures the amount of free morphine concentrations (ng/mL) in blood taken from 25 “acute” heroin overdose victims. The data are listed below:

67	110	45	96	80	103	223	86	58
112	52	165	91	189	85	150	120	293
153	390	47	158	77	216	526		

- Source: B Levine, D Green-Johnson, KA Moore, D Fowler, A Jenkins, (2001). Assessment of the acuteness of heroin deaths from the analysis of multiple blood specimens. Science and Justice 42(1). 17-20.

Heroin data (ordered)

```
heroin<-read.csv( 'DataSets/morphine.csv' )  
sort(heroin$Morphine)
```


```
## [1] 45 47 52 58 67 77 80 85 86 91 96  
## [12] 103 110 112 120 150 153 158 165 189 216 223  
## [23] 293 390 526
```

- Easy to see minimum and maximum
- What is a typical value?
- How variable are the data?
- Any clusters in the data?
- Any outliers?

Data visualization - dotplots

- A dotplot provides a simple 1-D representation of a quantitative variable.
 - Each dot represents one observation.
 - Provides instant visual summary of the data.

```
stripchart(heroin$Morphine, pch=16, vertical=FALSE)
```



Describing distributions

- The shape of a histogram can be described by a smooth curve that roughly follows the tops of the bars.
- This can be done by eye or by numerical algorithms on the computer.
- The shape often follows certain patterns:
 - symmetric (same right to left)
 - bell-shaped
 - skewed to the right or left
 - unimodal (one major peak)
 - bimodal (two major peaks)
 - spread (wide or narrow)

Summarizing quantitative data

- Describe the distribution of a collection of data
- Measures of centre or location
 - Means or medians
- Measures of spread or scale
 - Standard deviations, variances, ranges, and interquartile ranges

Measuring the centre

- Average or sample mean
- The mean is *“the sum of all the observations divided by the number of observations”*
- In R:

```
mean(heroin$Morphine)
```

```
## [1] 147.68
```


Measuring the centre

- The sample median is the middle value when the observations are ordered.
- Rule
 - If you have an odd number of observations, then the median is the middle value.
 - If you have an even number, the median is the average of the two middle observations.
- In R:

```
median(heroin$Morphine)
```

```
## [1] 110
```

Which one should we use?

- Mean
 - The “balance point” of the histogram
 - Affected by outliers (e.g., heroin data)
 - Most commonly used in testing and estimation
- Median
 - Divides the area of the histogram in half
 - Unaffected by a few outliers
- If the histogram is fairly symmetric, the mean and median will be similar.

Five number summary

- When the data are ordered
 - LQ is $\frac{1}{4}$ of the way in from the bottom
 - UQ is $\frac{1}{4}$ of the way from the top


```
summary(heroin$Morphine)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	45.0	80.0	110.0	147.7	165.0	526.0

Visualization: boxplots

- Boxplots provide another method for summarizing a quantitative variable.
 - Box contains central 50% of data (between LQ and UQ).
 - Median is line across box.
- Software packages differ in calculation of “whisker length”.

```
boxplot(heroin$Morphine, horizontal = TRUE)
```



Measures of Dispersion

- There are several ways to describe the variability or spread of a data set:
 - Range
 - Interquartile range
 - Standard deviation
 - Variance
- Standard deviation is the most common.

Ranges

- The Sample Range: difference between the largest and smallest observation.
 - Advantage: easy to calculate.
 - Disadvantages: highly sensitive; not useful for testing
- The Interquartile Range (IQR): difference between the first and third quartiles: $IQR = Q3 - Q1$.

```
range(heroin$Morphine)
```

```
## [1] 45 526
```

```
IQR(heroin$Morphine)
```

```
## [1] 85
```

Sample Standard Deviation

- Measures how far away, on average, values are from the center
- Same units as the data
- $s=0$ only when there is no spread, otherwise $s>0$.

```
sd(heroin$Morphine)
```

```
## [1] 113.3735
```

The Sample Variance

- The sample variance is just the square of the sample standard deviation. Its symbol is s^2 .
- “Average squared distance from the mean.”
- Units are the square of the data units.

```
var(heroin$Morphine)
```

```
## [1] 12853.56
```

```
sd(heroin$Morphine)^2
```

```
## [1] 12853.56
```


Summary - describing “spread”


- Range = maximum – minimum
- Not sensible – affected severely by outliers
- Interquartile range = UQ – LQ
- Describes the central half (50%) of the data
- Sample standard deviation (s.d.) = “square root of the averaged squared distance to the mean”

Estimation - reporting uncertainty

- Once we have gathered/generated/downloaded some data, we often want to estimate some quantity that is of interest:
 - e.g., the mean of the population from which our sample was taken
- We can do this by using the concept of the *confidence interval*
 - uses quantities calculated from the data to obtain population-level estimates
 - report the *point estimate* (i.e., a single number), along with an estimate of our uncertainty

The Normal Distribution

- Describes a wide variety of data very well
 - e.g., measurement error, height, body temperature, levels in blood samples (e.g. cholesterol)
- Characterized by the (population) mean, μ (defines centre) and the (population) SD σ (defines spread).
- 68% of distribution is between $-\sigma$ and $+\sigma$, 95% between -2σ and $+2\sigma$ (actually 1.96) and 99.7% between -3σ and $+3\sigma$.



Sampling Distributions

- Start with population with a certain distribution (e.g., normal, but doesn't have to be).
- True mean μ and SD σ (parameters describing distribution)
- Take sample of size n from normal population
 - Get sample y_1, \dots, y_n
 - Calculate mean, \bar{y} and sample SD s (estimates for parameters).
- Close(ish) to μ and σ , especially if n is large.
- What if we drew a different sample? Wouldn't get same \bar{y} and s , since sample would be different
- Sampling variability: how does taking different samples affect our estimates of μ , and σ ?

Meta-experiment

- Think about repeating experiment many times (taking one sample of n is one repetition).
- Think of the probability distribution of all the possible outcomes for the estimates \bar{y} , s .
- This is called a sampling distribution for the parameter estimates (here \bar{y} , s).
- Usually can't actually generate the sampling distribution: can only do sampling once
- Sampling distribution represents a conceptual *meta-experiment*.

Visualising the “meta-experiment”

- Go to the website below, click on “Chapters” (next to “Home” at the top left of the screen) and select “Probability Distributions” and then “Central Limit Theorem”:

<http://students.brown.edu/seeing-theory/index.html>

- Set the parameters to $\alpha = 5$ and $\beta = 5$, with a sample size of 10 (technically not a normal distribution, but close enough for our purposes).
- Generate draws from this distribution, and observe the shape of the resulting distribution of the means (each “draw” generates 10 observations, and takes the mean: the distribution of means is plotted in the lower graph).
- The shape should be suspiciously familiar...

Visualising the “meta-experiment”

≡

Chapter 3: Probability Distributions

Central Limit Theorem

The Central Limit Theorem (CLT) states that the sample mean of a sufficiently large number of i.i.d. random variables is approximately normally distributed. The larger the sample, the better the approximation.

Change the parameters α and β to change the distribution from which to sample.

$\alpha = 5.00$


$\beta = 5.00$

Choose the sample size and how many sample means should be computed (draw number), then press "Sample." Check the box to display the true distribution of the sample mean.

Sample size = 10

Draws = 50

☒ Theoretical



<http://students.brown.edu/seeing-theory/index.html>

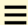
Sampling Distribution: Normal Mean

- *What does it look like?*
 - it looks like a normal curve
 - it has same mean as the original distribution: μ .
 - it has less variability: the SD is smaller.
 - This is because some of the “noise” in the individual observations is smoothed out in the \bar{y} 's.
 - new distribution has $SD = \sigma/\sqrt{n}$.

Central Limit Theorem

- Even if Y doesn't have a normal distribution, the sampling distribution of \bar{y} will be close to normal, when n is large.
 - It also has mean μ and SD equal to σ/\sqrt{n}
- What does "large" mean in this theorem?
- It depends how close to normal the distribution of Y is.
 - If Y is normally distributed, then $n=1$ is big enough.
 - If Y is not normally distributed, but is fairly symmetric and fairly bell-shaped, then n doesn't have to be too large (e.g., $n=10$).
 - If the distribution of Y is really weird-looking, then n may have to be really big.


Visualising the Central Limit Theorem


 Chapter 3: Probability Distributions

Central Limit Theorem


The Central Limit Theorem (CLT) states that the sample mean of a sufficiently large number of i.i.d. random variables is approximately normally distributed. The larger the sample, the better the approximation.


Change the parameters α and β to change the distribution from which to sample.

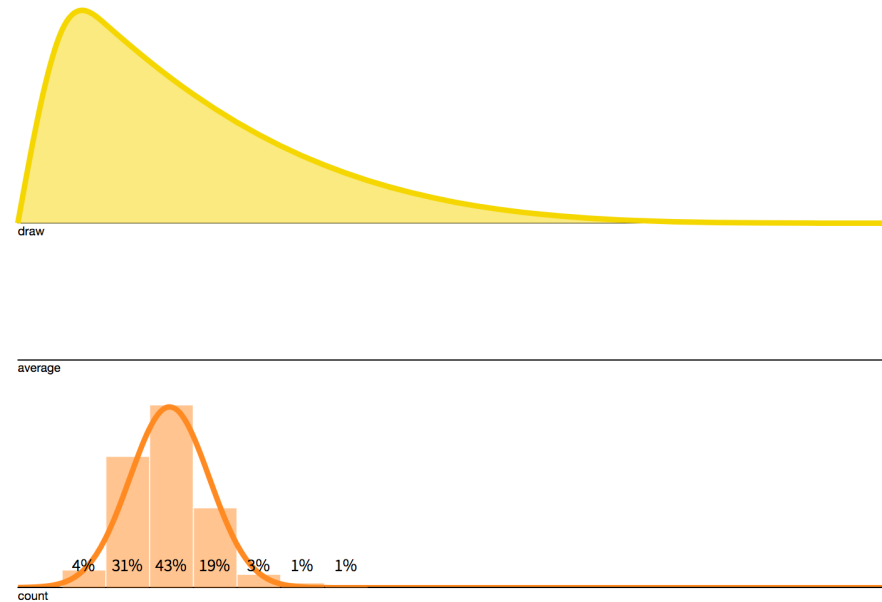
$\alpha = 1.05$ 

$\beta = 5.00$ 

Choose the sample size and how many sample means should be computed (draw number), then press "Sample." Check the box to display the true distribution of the sample mean.

Sample size = 10 

Draws = 50 



<http://students.brown.edu/seeing-theory/index.html>

Estimation


- Usually do not know the true μ and σ .
- Want to estimate μ with \bar{y} and also have some idea how good our estimate is.
- Want to make an interval around \bar{y} that we are, say, 95% certain contains the true value of μ .
 - called a 95% confidence interval
- In general we want to find a “ $1-\alpha$ ” confidence interval
 - For a 95% CI, have $\alpha = 0.05$; for a 90% CI have $\alpha = 0.10$, for a 99% CI have $\alpha = 0.01$, etc

Estimation (continued)

- We do the best we can: must estimate σ with s .
- Define the *standard error of the mean* as $SE = s/\sqrt{n}$.
- SE is an estimate of the true standard deviation of the mean, which is σ/\sqrt{n} .
 - Will use the estimate SE in place of the true value .
- There is a price to pay for not using the true value: can't use normal distribution for our calculations.
 - Using estimate for introduces extra uncertainty.
 - Makes the confidence interval wider.

The t distribution

- Use a distribution with a little more variability than the normal distribution: the *t-distribution*
 - shape similar to the standard normal $N(0,1)$ but kind of squashed so that it has a wider spread
 - always centered at mean 0, with amount of squashing determined by the “degrees of freedom” (df).



Example: confidence interval

- Suppose we have the following data on daily calcium intake (mg) from a study of women's bone health:

```
calcium<-read.csv('DataSets/calcium.csv')  
head(calcium)
```

```
##   ID Calcium  
## 1  1     808  
## 2  2     882  
## 3  3    1062  
## 4  4     970  
## 5  5     909  
## 6  6     802
```

- What is our estimate of the average daily calcium intake?

Example: confidence interval

```
t.test(calcium$Calcium)

##
## One Sample t-test
##
## data:  calcium$Calcium
## t = 13.361, df = 37, p-value = 9.895e-16
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##    785.5995 1066.4532
## sample estimates:
## mean of x
##  926.0263
```

Estimate: 926mg, 95% CI: (786mg, 1066mg)

Visualising Confidence Intervals

Confidence Interval

In contrast to point estimators, confidence intervals estimate a parameter by specifying a range of possible values. Such an interval is associated with a confidence level, which is the probability that the procedure used to generate the interval will produce an interval containing the true parameter.

Choose a probability distribution to sample from.



Normal

Choose a sample size (n) and confidence level ($1 - \alpha$).

$n = 10$

$1 - \alpha = 0.95$

Start sampling to generate confidence intervals.



<http://students.brown.edu/seeing-theory/index.html>

Standard Error vs Standard Deviation

- Sample SD s :
 - measures amount of variability in data set
 - as n gets larger, SD approaches true SD ; does not tend to get smaller or larger;
 - used for making predictions about single data point.
- Sample SE = s/\sqrt{n} :
 - measures amount of uncertainty in estimate of mean
 - gets smaller as n gets larger (random sample)
 - used for making predictions about mean.

Summary: estimation of μ

- Estimate the population mean by the sample mean, \bar{x} .
 - There is still uncertainty about the true value of μ after we calculate \bar{x} .
 - construct a confidence interval for μ based on the observed data.
- A confidence interval is a statement about a population parameter based on the observed data, **NOT** a statement about the sample.
- The width of the confidence interval increases as:
 - the level of confidence increases (i.e., 99% *versus* 95%)
 - the sample size (i.e., n) decreases
 - the variability (i.e., s) increases

Two independent samples

- Sometimes we have two populations (groups) we want to compare, based on a continuous measurement.
- Examples:
 - cases (treatment) and controls
 - drug and placebo
 - patients taking two comparable drugs
 - males and females
 - two different genetic strains, or mutant/wildtype
- Population 1: $N(\mu_1, \sigma_1)$, sample $n_1, \bar{y}_1, s_1, SE_1 = s_1/\sqrt{n_1}$
- Population 2: $N(\mu_2, \sigma_2)$, sample $n_2, \bar{y}_2, s_2, SE_2 = s_2/\sqrt{n_2}$

Comparing μ_1 and μ_2

- Typical question: is μ_1 larger than μ_2 ?
 - What is the difference in the means of the two populations, $\mu_1 - \mu_2$?
 - We want to estimate $\mu_1 - \mu_2$ and get a Confidence Interval.
- What to do:
 - take random samples from each population, sizes n_1 and n_2 .
 - calculate the mean from each sample
 - $\bar{y}_1 - \bar{y}_2$ is an estimate for $\mu_1 - \mu_2$
 - need a SE to find the width of a Confidence Interval

SE for difference between two means

- How to calculate SE for the estimate $\bar{y}_1 - \bar{y}_2$?
 - two ways: unpooled and pooled.
 - when $n_1 = n_2$ these are equivalent.
- Use Unpooled method most of the time.
- Use Pooled method when you think $\sigma_1 = \sigma_2$.
- Both are based on some combination of s_1 and s_2

Example: comparing two means

- Data from trial involving feeding day-old male chicks “high lysine corn”. Weight gain recorded (Cromwell, et al. Poultry Sci. 1968).


```
corn<-read.csv( 'DataSets/corn.csv' )  
head(corn)
```

```
##      Group Weight  
## 1 Control    380  
## 2 Control    321  
## 3 Control    366  
## 4 Control    356  
## 5 Control    283  
## 6 Control    349
```

- Is there a significance difference in weight gain between the groups?

Example: comparing two means

```
## Generate boxplots for each group  
boxplot(Weight ~ Group, data=corn)
```



Comparing two means: 95% CI

```
## Confidence interval for difference between means
```

```
t.test(Weight ~ Group, data=corn)
```

```
##
```

```
## Welch Two Sample t-test
```

```
##
```

```
## data: Weight by Group
```

```
## t = -2.469, df = 36.915, p-value = 0.0183
```

```
## alternative hypothesis: true difference in means is not equal to 0
```

```
## 95 percent confidence interval:
```

```
## -66.729186 -6.570814
```

```
## sample estimates:
```

```
## mean in group Control    mean in group Exper
```

```
##           366.30           402.95
```

95% confidence interval for difference: (-66.73, -6.57)

Comparing two means: 99% CI

```
## Confidence interval for difference between means  
t.test(Weight ~ Group, data=corn, conf.level = 0.99)
```

```
##  
## Welch Two Sample t-test  
##  
## data: Weight by Group  
## t = -2.469, df = 36.915, p-value = 0.0183  
## alternative hypothesis: true difference in means is not equal to 0  
## 99 percent confidence interval:  
## -76.962594 3.662594  
## sample estimates:  
## mean in group Control mean in group Exper  
## 366.30 402.95
```

99% confidence interval for difference: (-76.96, 3.66)

Interpretation of Confidence Intervals

- If we construct all our confidence intervals with this procedure, in the long run
 - we will be right 95% = $1 - \alpha$ of the time
 - we will be wrong 5% = α of the time
- α represents our “willingness” to be wrong
- In reality we cannot repeat the experiment, and we do not know μ .
- Therefore we NEVER KNOW for any one experiment whether we are right or wrong

Hypothesis testing: The t-test

- Can compare population means (e.g., μ_1 and μ_2) using confidence intervals.
- Often interested in whether means are the same ($\mu_1 = \mu_2$) or different ($\mu_1 \neq \mu_2$).
- Can also accomplish this comparison using Hypothesis Testing.

Null and alternative hypotheses

- Two populations, means μ_1 and μ_2 .
 - Null hypothesis, $H_0 : \mu_1 = \mu_2$
 - Alternative hypothesis, $H_A : \mu_1 \neq \mu_2$
- Idea is to determine whether the data contains enough evidence to make H_0 seem unlikely.
- Example: does studying improve exam performance?
 - Two populations: the exam scores for those who didn't study (population mean 1) and the exam scores for those who did (population mean 2).
 - H_0 : Studying has no effect on exam score ($\mu_1 = \mu_2$).
 - H_A : Studying has some effect on exam score ($\mu_1 \neq \mu_2$).

The t statistic

- Testing the null hypothesis:

$$H_0 : \mu_1 = \mu_2 \iff \mu_1 - \mu_2 = 0$$

- Against the alternative hypothesis

$$H_A : \mu_1 \neq \mu_2 \iff \mu_1 - \mu_2 \neq 0$$

- Use t statistic (assume H_0 is true):

$$t_s = \frac{(\bar{y}_1 - \bar{y}_2) - (\mu_1 - \mu_2)}{SE(\bar{y}_1 - \bar{y}_2)} = \frac{(\bar{y}_1 - \bar{y}_2)}{SE(\bar{y}_1 - \bar{y}_2)}$$

Example – high lysine corn data

```
t.test(Weight ~ Group, data=corn)
```

```
##  
## Welch Two Sample t-test  
##  
## data: Weight by Group  
## t = -2.469, df = 36.915, p-value = 0.0183  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -66.729186 -6.570814  
## sample estimates:  
## mean in group Control    mean in group Exper  
##           366.30           402.95
```

Example – high lysine corn data

```
aggregate(corn$Weight, list(corn$Group), mean)
```

```
##   Group.1      x  
## 1 Control 366.30  
## 2  Exper 402.95
```

```
aggregate(corn$Weight, list(corn$Group), sd)
```

```
##   Group.1      x  
## 1 Control 50.80520  
## 2  Exper 42.72861
```

$$t_s = \frac{(\bar{y}_1 - \bar{y}_2)}{SE(\bar{y}_1 - \bar{y}_2)} = \frac{366.30 - 402.95}{\sqrt{\frac{50.81^2}{20} + \frac{42.73^2}{20}}} = -2.469$$

The P-value

- The P-value of the test is the area of the t distribution in the “double tails”, beyond $-t_s$ and $+t_s$.
 - Measure of compatibility between data and H_0 .
- The P-value for a hypothesis test is the probability, computed under the condition that the null hypothesis is true, of the test statistic being at least as extreme as the value of the test statistic that was actually obtained.
- Basically we are asking: “what is the probability of observing these data, if the null hypothesis is true?”


Example: high lysine corn

```
t.test(Weight ~ Group, data=corn)
```

```
##  
## Welch Two Sample t-test  
##  
## data: Weight by Group  
## t = -2.469, df = 36.915, p-value = 0.0183  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -66.729186 -6.570814  
## sample estimates:  
## mean in group Control mean in group Exper  
## 366.30 402.95
```

- So we would only expect to observe a t statistic as large as this ($t_s = -2.47$) in 1.83% of samples when H_0 is true.
- Based on this we are more likely to believe H_1 .

Visualising P-values (old site)



p-Values

The p-value is the probability that a statistical summary, such as mean, would be the same as or more extreme than an observed result given a probability distribution for the statistical summary.

1. Choose a probability distribution for the statistical summary.

Choose ▾ studentt (37)

2. Decide what type of p-value to compute.

☐ One-Tailed ☒ Two-Tailed

3. Type in an observed result to visualize the p-value.

2.469

$p = 0.0183$

4. Switch input from observation to a p-value and visualization computes the critical value....

<http://students.brown.edu/seeing-theory/index.html>

Drawing conclusions from a t test

- How to interpret P-values?
 - Significance level: $\alpha = 0.1, 0.05, 0.01$.
 - Value of α is chosen by whoever is making the decision.
- If P-value is less than or equal to α , H_0 is rejected.
- For the corn example, the P-value was 0.0183.
 - Reject H_0 at $\alpha = 0.1$ or $\alpha = 0.05$, but not at $\alpha = 0.01$.

Type I and Type II errors

- Type I error: reject H_0 when H_0 is in fact true.
- Type II error: fail to reject H_0 when H_0 is false.

	True situation	
	H_0 true	H_0 false
Do not reject H_0	Correct	Type II error
Reject H_0	Type I error	Correct

Interpreting

- How do you know which value of α to use?
 - α is the probability of rejecting H_0 when H_0 is true.
 - i.e., $P(\text{Type I error}) = \alpha$
- If you conducted a large number of hypothesis tests in which the null hypothesis were true, the expected proportion of tests in which H_0 would be rejected is equal to α .

Significant versus *important*

- Difference between μ_1 and μ_2 is said to be *significant* if the P-value is less than or equal to α . Otherwise the difference is said to be *non-significant*.
- **BUT**, just because a difference is significant, it doesn't make it important. Can't use the P-value to decide whether a significant difference is important. Have to think about importance in the context of the problem.

Example: Serum LD

- Serum LD: enzyme showing elevated levels after heart damage

	Males	Females
n	270	264
\bar{y}	60	57
s	11	10

$$t_s = \frac{(\bar{y}_1 - \bar{y}_2)}{SE(\bar{y}_1 - \bar{y}_2)} = \frac{60 - 57}{\sqrt{\frac{11^2}{270} + \frac{10^2}{264}}} = 3.29 \implies p \approx 0.001$$

- So the difference is statistically significant, but a difference of 3 between males and females may not be important.

Example: Serum LD

